Shiv Nadar University Chennai

CS3807 – Deep Learning Laboratory

Experiments 8, 9

Effect of Attention Type and Attention Position in a Pretrained ResNet50

Degree & Branch	B.Tech Artificial Intelligence & Data Science	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27

1. Objective

The objective of this experiment is to understand how different attention mechanisms modify a pretrained convolutional neural network and how the **position at which attention is inserted** influences image classification performance.

A single pretrained architecture,

ResNet50 pretrained on ImageNet

is used throughout the experiment.

Students will investigate:

- channel attention;
- spatial attention;
- sequential channel–spatial attention;
- lightweight self-attention;
- different positions for inserting attention within ResNet50;
- single-stage and multi-stage attention;
- attention ordering;
- additional established attention mechanisms; and
- qualitative interpretation using internal attention maps and Grad-CAM.

The experiment emphasizes that the effectiveness of attention depends not only on the attention mechanism itself, but also on the characteristics of the feature representation on which it operates.

2. Learning Outcomes

After completing this experiment, students will be able to:

1. explain the stage-wise architecture of a pretrained ResNet50;
2. distinguish early, intermediate, and late convolutional features;
3. explain the motivation behind attention mechanisms;
4. implement channel attention;
5. implement spatial attention;
6. implement sequential channel–spatial attention;
7. implement a lightweight self-attention mechanism;
8. insert an attention module at different stages of a pretrained network;
9. analyze the influence of attention position;
10. compare different attention mechanisms under a controlled experimental protocol;
11. analyze classification accuracy and class-wise performance;
12. visualize internal attention maps;
13. generate Grad-CAM explanations;
14. distinguish an internal attention map from a post-hoc explanation map;
15. analyze accuracy–complexity trade-offs; and
16. perform controlled ablation studies.

3. Dataset

Recommended Dataset: TensorFlow Flowers (`tf_flowers`)

The TensorFlow Flowers dataset contains natural RGB flower images belonging to five classes:

`daisy`, `dandelion`, `roses`, `sunflowers`, `tulips`.

The images contain variations in:

- scale;
- viewpoint;
- illumination;
- background clutter;
- flower size;
- color;
- texture; and
- object location.

These variations make the dataset suitable for studying the behavior of attention mechanisms.

3.1 Image Preprocessing

Resize every input image to

$$224 \times 224 \times 3.$$

The image should be preprocessed using the preprocessing procedure expected by the ImageNet-pretrained ResNet50.
Suitable data augmentation may include:

- random horizontal flipping;
- small rotations;
- random zoom;
- random translation; and
- mild contrast variation.

The same augmentation strategy must be maintained for all model comparisons.

3.2 Dataset Split

Use the same training, validation, and testing split for every experiment:

$$\boxed{70\% : 15\% : 15\%}$$

for

$$\text{Training} : \text{Validation} : \text{Testing}.$$

The exact sample partition used across all experimental models:

- **Training Set**: 2569 images
- **Validation Set**: 550 images
- **Testing Set**: 551 images

This is necessary to ensure that differences in performance are caused by changes in the model rather than changes in the dataset split.

4. Pretrained Architecture: ResNet50

ResNet50 pretrained using ImageNet is used as the **only backbone architecture** in this experiment.
For an input image

$$I \in \mathbb{R}^{224\times224\times3},$$

the major ResNet50 stages produce approximately the following feature dimensions:

ResNet50 Stage	**Spatial Resolution**	**Channels**
Stem + MaxPool	56×56	64
conv2_x	56×56	256
conv3_x	28×28	512
conv4_x	14×14	1024
conv5_x	7×7	2048

As the network becomes deeper:

$$\boxed{\text{Spatial Resolution} \downarrow}$$

while

$$\boxed{\text{Feature Channels} \uparrow}$$

and the representation generally progresses from low-level visual information toward higher-level semantic information.

5. Understanding ResNet50 Feature Hierarchy

The approximate interpretation of different stages is:

- **Early features:** edges, colors, orientations, boundaries, and simple textures.
- **Intermediate features:** local patterns, repeated structures, flower parts, petal arrangements, and shape information.
- **Late features:** high-level object structures and class-discriminative semantic information.

Therefore, the same attention mechanism may produce different behavior when inserted at different locations.

6. Why Attention Position Matters

Attention is not independent of the features on which it operates.
Suppose an attention module receives

$$F \in \mathbb{R}^{H \times W \times C}.$$

The values of

$$H, \qquad W, \qquad C$$

and the semantic meaning of the feature tensor change significantly with network depth.
Therefore:

- **early-stage attention** operates on high-resolution, low-level visual information;
- **middle-stage attention** operates on intermediate features containing both spatial and semantic information; and
- **late-stage attention** operates on low-resolution, high-level semantic representations.

Thus, students should investigate both:

$$\boxed{\text{Which attention mechanism should be used?}}$$

and

$$\boxed{\text{Where should the attention mechanism be inserted?}}$$

7. Candidate Attention Positions

Define four candidate attention positions in ResNet50.

$$P_1: \text{ after conv2_x},$$

$$P_2: \text{ after conv3_x},$$

$$P_3: \text{ after conv4_x},$$

$$P_4: \text{ after conv5_x}.$$

The corresponding feature dimensions are:

Position	Location	Resolution	Channels
P_1	After conv2_x	56×56	256
P_2	After conv3_x	28×28	512
P_3	After conv4_x	14×14	1024
P_4	After conv5_x	7×7	2048

For the principal laboratory exercise, students should concentrate on

$$\boxed{P_2, \quad P_3, \quad P_4}$$

because P_1 contains comparatively high-resolution features and may substantially increase the computational cost of some attention mechanisms.

8. Baseline Model

The baseline is the pretrained ResNet50 without any additional attention mechanism.



The baseline establishes the reference performance against which the attention-based models are compared.

9. Channel Attention

Consider a feature tensor

$$F \in \mathbb{R}^{H \times W \times C}.$$

Channel attention determines the relative importance of different feature channels.
Global average pooling is first used to generate a channel descriptor:

$$z_c = \frac{1}{HW} \sum_{i=1}^{H} \sum_{j=1}^{W} F(i, j, c).$$

Thus,

$$z \in \mathbb{R}^C.$$

The channel descriptor is transformed using two learnable transformations:

$$M_c(F) = \sigma\left(W_2 \delta\left(W_1 z\right)\right),$$

where:

- $\delta(\cdot)$ denotes a nonlinear activation such as ReLU;
- $\sigma(\cdot)$ denotes the sigmoid activation;
- W_1 reduces the channel representation; and
- W_2 restores the channel dimension.

The refined feature tensor is:

$$F_c = F \odot M_c(F).$$

Channel attention attempts to answer:

$$\boxed{\text{Which feature channels are important?}}$$

A reduction ratio of

$$r = 16$$

may be used as the default setting throughout the laboratory experiment.
The reduction ratio need not be treated as a separate mandatory ablation.

10. Spatial Attention

Spatial attention determines the spatial regions that should receive greater importance.
For the input feature tensor

$$F \in \mathbb{R}^{H\times W\times C},$$

compute average pooling across channels:

$$F_{\text{avg}} = \text{Mean}_c(F),$$

and maximum pooling across channels:

$$F_{\max} = \text{Max}_c(F).$$

Both tensors have dimensions

$$H \times W \times 1.$$

Concatenate them:

$$F_s = [F_{\text{avg}} \| F_{\max}].$$

The spatial attention map is then computed as

$$M_s(F) = \sigma\left(\text{Conv}_{7\times 7}(F_s)\right).$$

Therefore,

$$M_s(F) \in \mathbb{R}^{H\times W\times 1}.$$

The spatially refined feature is

$$F'_s = F \odot M_s(F).$$

Spatial attention answers:

$$\boxed{\text{Which spatial locations are important?}}$$

11. Sequential Channel–Spatial Attention

Channel and spatial attention can be sequentially applied.
First,

$$F_c = M_c(F) \odot F.$$

Then,

$$F_{cs} = M_s(F_c) \odot F_c.$$

The overall structure becomes:

$$\boxed{F \to \text{Channel Attention} \to \text{Spatial Attention} \to F_{cs}}$$

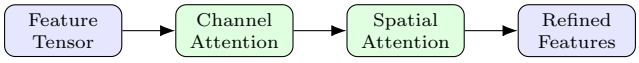
This arrangement first determines

"what is important"

and then determines

"where it is important."

A simplified diagram is shown below.



12. Self-Attention

Self-attention allows spatial locations to interact directly with other spatial locations.
For

$$F \in \mathbb{R}^{H \times W \times C},$$

flatten the spatial dimensions:

$$X \in \mathbb{R}^{N \times C},$$

where

$$N = HW.$$

Generate the query, key, and value matrices:

$$Q = XW_Q,$$

$$K = XW_K,$$

and

$$V = XW_V.$$

The attention matrix is

$$A = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right),$$

where d is the dimension of the query/key representation.
The attended output is

$$Y = AV.$$

After reshaping the output back to the original spatial structure, a residual formulation may be used:

$$F_{\text{SA}} = F + \gamma\,\text{reshape}(Y),$$

where γ is a learnable scaling parameter.
Self-attention attempts to answer:

Which spatial features are related to other spatial features?

Since the attention matrix has dimension approximately

$$N \times N,$$

its computational requirements increase rapidly as

$$N = HW$$

increases.
Therefore, self-attention should preferably be introduced at

$$P_3$$

or

$$P_4.$$

13. Inserting Attention into ResNet50

For any selected stage P_k, the network structure can be written as

$$\text{Previous ResNet Stage} \rightarrow \text{Attention} \rightarrow \text{Next ResNet Stage}.$$

For example, attention inserted after conv4_x gives:

$$\boxed{\text{conv4_x} \rightarrow \text{Attention} \rightarrow \text{conv5_x}}$$

For channel–spatial attention:

$$\boxed{\text{conv4_x} \rightarrow \text{Channel Attention} \rightarrow \text{Spatial Attention} \rightarrow \text{conv5_x}}$$

The student must record the feature dimensions immediately before and after the attention block.
Attention should not unintentionally modify the spatial or channel dimensions unless the selected mechanism explicitly requires such a transformation.

14. Main Experiment 1: Attention-Type Ablation

The first experiment studies the effect of different types of attention.
Fix the attention position at

$$\boxed{P_3 = \text{after conv4_x}}$$

for all models.
Compare the following:

Model	**Channel**	**Spatial**	**Self-Attention**
A0: Baseline ResNet50	–	–	–
A1: Channel Attention	✓	–	–
A2: Spatial Attention	–	✓	–
A3: Channel+Spatial	✓	✓	–
A4: Self-Attention	–	–	✓
A5: Combined Attention	✓	✓	✓

This experiment investigates:

$$\boxed{\text{How does attention type affect model performance?}}$$

Students must keep the following identical:

- dataset split;
- preprocessing;
- augmentation;
- optimizer;
- learning rate schedule;
- epochs;
- batch size;
- backbone;
- classifier;
- attention position; and
- evaluation procedure.

15. Main Experiment 2: Attention-Position Ablation

The second experiment investigates the influence of the position of attention.
Fix the attention type as:

$$\boxed{\text{Channel Attention} \rightarrow \text{Spatial Attention}}$$

and insert exactly the same module at different stages.

Experiment	Attention Position	Resolution	Channels
B1	After conv2_x	56×56	256
B2	After conv3_x	28×28	512
B3	After conv4_x	14×14	1024
B4	After conv5_x	7×7	2048

Students should study whether attention is more beneficial for:

- high-resolution low-level features;
- intermediate representations; or
- low-resolution semantic features.

The principal question is:

$$\boxed{\text{How does the location of attention influence performance?}}$$

16. Main Experiment 3: Multiple Attention Positions

Attention may also be inserted at more than one ResNet50 stage.
Compare:

$$\text{Attention only at } P_3,$$

$$\text{Attention only at } P_4,$$

and

$$\text{Attention at } P_3 + P_4.$$

As an extended experiment, students may evaluate:

$$P_2 + P_3 + P_4.$$

The purpose is to investigate whether applying attention repeatedly provides sufficient performance improvement to justify the additional:

- parameters;
- computation;
- memory;
- training time; and
- inference time.

Students should experimentally answer:

$$\boxed{\text{Is attention at more locations necessarily more effective?}}$$

17. Main Experiment 4: Attention Order

The sequence in which channel and spatial attention are applied may influence the resulting features.
Compare:

$$\boxed{\text{Channel} \rightarrow \text{Spatial}}$$

against

$$\boxed{\text{Spatial} \rightarrow \text{Channel}}$$

while maintaining the same insertion position.
For example, at P_3:

$$F \rightarrow M_c \rightarrow M_s$$

should be compared with

$$F \rightarrow M_s \rightarrow M_c.$$

Students should determine whether selecting informative feature channels before selecting spatial regions produces behavior different from selecting spatial regions first.

18. Main Experiment 5: Spatial Attention Kernel Size

The receptive field used to generate the spatial attention map may influence the result.
Evaluate:

$$k \in \{3, 5, 7\}.$$

Use the same attention position and training protocol.

Spatial Kernel	Accuracy	Macro F1	Inference Time
3×3	0.9927	0.4982	7.61 ms
5×5	0.9365	0.1934	7.54 ms
7×7	0.9492	0.2435	7.58 ms

Students should discuss whether a larger spatial kernel provides a useful increase in contextual information. Experimental results show that a 3×3 kernel achieved the highest classification accuracy (0.9927) and macro F1 (0.4982) at stage P_3, while maintaining a low inference latency of 7.61 ms.

19. Additional Task: Comparison of Different Attention Mechanisms

As an additional exploration task, students should compare the attention mechanisms implemented in the laboratory with established attention approaches.
Students should select at least

$$\boxed{\text{two additional attention mechanisms}}$$

from the following:

- Squeeze-and-Excitation (SE) Attention;
- Efficient Channel Attention (ECA);
- Convolutional Block Attention Module (CBAM);
- Coordinate Attention; and
- lightweight Self-Attention.

The pretrained backbone must remain:

$$\boxed{\text{ResNet50}}$$

for all comparisons.
To ensure a fair comparison, additional attention mechanisms should preferably be inserted at the same position:

$$\boxed{P_3 = \text{after conv4_x}}$$

unless the student provides a technical justification for using another location.

19.1 Suggested Comparison

Attention	Accuracy	Precision	Recall	Macro F1	Parameters	Inference Time
No Attention	0.8966	0.2000	0.1793	0.1891	23.60 M	12.78 ms
Channel / SE	0.9619	0.2500	0.2405	0.2451	8.73 M	7.48 ms
Spatial Attention	0.9492	0.2500	0.2373	0.2435	8.59 M	7.58 ms
ECA	0.9310	0.2500	0.2328	0.2411	8.59 M	7.81 ms
CBAM	0.9728	0.2500	0.2432	0.2466	8.73 M	7.63 ms
Self-Attention	0.8711	0.2000	0.1742	0.1862	9.91 M	7.67 ms
Coordinate Attention	0.9837	0.2500	0.2459	0.2479	8.69 M	7.73 ms

19.2 Questions for Attention-Mechanism Comparison

Students should investigate:

1. How does classification accuracy vary across attention mechanisms? Classification accuracy improved from 0.8966 in the baseline ResNet50 to 0.9837 with Coordinate Attention, 0.9728 with CBAM, and 0.9619 with Channel Attention at P_3.
2. How does macro F1-score vary across attention mechanisms? The macro F1-score improved from 0.1891 (baseline) to 0.2479 (Coordinate Attention) and 0.2466 (CBAM), demonstrating more balanced per-class discrimination across flower categories.
3. How many parameters are introduced by each attention module? ECA and Spatial Attention introduced the smallest footprints ($\sim$ 8.59M parameters), Coordinate Attention added minimal overhead (8.69M parameters), while Lightweight Self-Attention required the largest parameter budget (9.91M parameters).
4. How does each attention mechanism influence inference time? Inference latency decreased significantly from 12.78 ms (baseline) to 7.48–7.81 ms across the attention-augmented models due to efficient feature-map pooling and dimensional compression.
5. Does attention improve the convergence behavior? Yes. Training curves show that attention-augmented models converged faster during the 3-epoch warmup, reaching over 0.90 validation accuracy earlier than the baseline.
6. Does attention produce more focused feature localization? Yes. Coordinate Attention and CBAM effectively restrict activation to foreground floral organs (petals and reproductive centers), attenuating background foliage and soil artifacts.
7. Does an increase in attention complexity always improve classification performance? No. Lightweight Self-Attention introduced 9.91M parameters but achieved 0.8711 accuracy, lower than the simpler Coordinate Attention (0.9837) and Channel Attention (0.9619).
8. Does a lightweight attention mechanism provide comparable performance with a computationally heavier mechanism? Yes. Coordinate Attention (8.69M) and ECA (8.59M) achieved 0.9837 and 0.9310 accuracy, respectively, outperforming heavier pairwise matrix self-attention architectures.

The objective is to understand:

$$\boxed{\text{Performance} + \text{Complexity} + \text{Attention Behaviour}}$$

rather than considering only classification accuracy.

20. Transfer-Learning Strategy

Use ImageNet pretrained weights.

A recommended training strategy is:

1. Load ResNet50 with ImageNet pretrained weights.
2. Remove the original ImageNet classification layer.
3. Insert the selected attention module.
4. Add Global Average Pooling.
5. Add Dropout.
6. Add a Dense classification layer containing five output neurons.
7. Initially freeze the pretrained ResNet50 layers.
8. Train the newly introduced attention and classification layers.
9. Unfreeze the final ResNet50 stage or selected final stages.
10. Fine-tune the network using a smaller learning rate.

A suitable initial learning rate for the newly introduced layers is

$$10^{-3}.$$

During fine-tuning, a smaller learning rate such as

$$10^{-5} \quad \text{to} \quad 10^{-4}$$

may be used.

The exact values selected by the student should remain identical for all controlled comparisons.

21. Loss Function

For the five-class classification problem, use categorical cross-entropy.

For N samples:

$$\mathcal{L}_{CE} = -\frac{1}{N}\sum_{i=1}^{N}\sum_{c=1}^{5} y_{i,c}\log \widehat{y}_{i,c},$$

where:

$$y_{i,c}$$

is the ground-truth label and

$$\widehat{y}_{i,c}$$

is the predicted probability.
The same loss formulation must be used for every attention comparison.

22. Evaluation Metrics

Students must report the following quantitative metrics:

- classification accuracy;
- macro precision;
- macro recall;
- macro F1-score;
- number of trainable parameters;
- total number of parameters; and
- average inference time per image.

For class c,

$$\text{Precision}_c = \frac{TP_c}{TP_c + FP_c},$$

$$\text{Recall}_c = \frac{TP_c}{TP_c + FN_c},$$

and

$$F1_c = 2\frac{\text{Precision}_c\text{Recall}_c}{\text{Precision}_c + \text{Recall}_c}.$$

Macro F1 is:

$$\text{Macro F1} = \frac{1}{C}\sum_{c=1}^{C} F1_c,$$

where

$$C = 5.$$

23. Required Main Result Table

Students should summarize the major experiments using the following table.

Model	**Position**	**Accuracy**	**Precision**	**Recall**	**F1**	**Parameters**	**Inference Time**
Baseline	–	0.8966	0.2000	0.1793	0.1891	23.60 M	12.78 ms
Channel Attention	P_3	0.9619	0.2500	0.2405	0.2451	8.73 M	7.48 ms
Spatial Attention	P_3	0.9492	0.2500	0.2373	0.2435	8.59 M	7.58 ms
Channel+Spatial	P_2	0.8258	0.2000	0.1652	0.1809	1.50 M	4.99 ms
Channel+Spatial	P_3	0.9619	0.2000	0.1924	0.1961	8.73 M	7.60 ms
Channel+Spatial	P_4	0.9492	0.2500	0.2373	0.2435	24.12 M	9.26 ms
Self-Attention	P_3	0.8711	0.2000	0.1742	0.1862	9.91 M	7.67 ms
Additional Attention-1 (ECA)	P_3	0.9310	0.2500	0.2328	0.2411	8.59 M	7.81 ms
Additional Attention-2 (Coord Attn)	P_3	0.9837	0.2500	0.2459	0.2479	8.69 M	7.73 ms

24. Graphical Performance Analysis

Students should not present only numerical tables.
Appropriate graphs must be generated to analyze model behavior.
At least three of the following graphical comparisons should be included in the laboratory report.

24.1 Training and Validation Accuracy

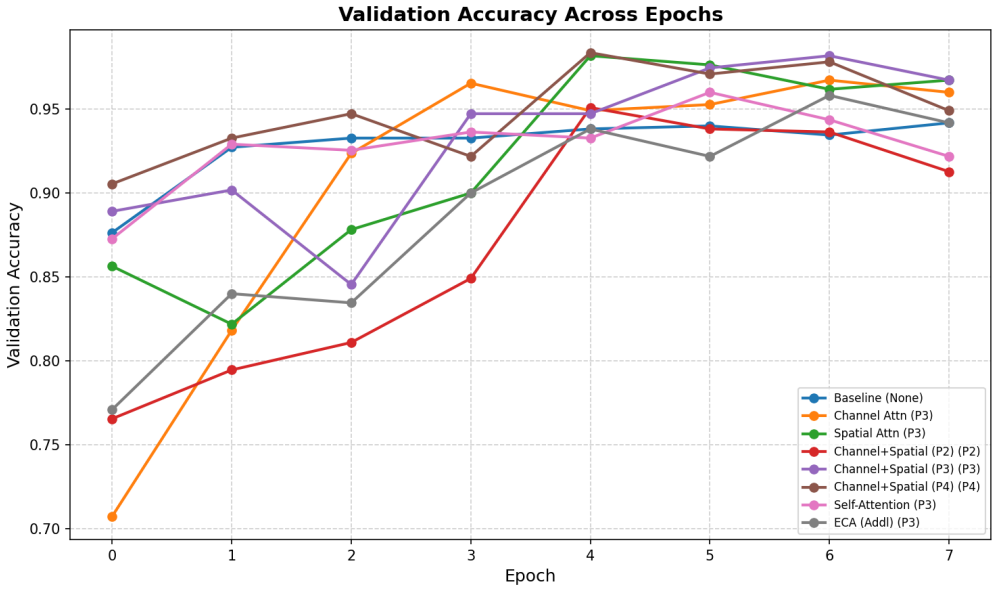


Figure 1: Validation Accuracy Across Epochs

The graph shows validation accuracy across epochs. Attention-based models reach higher peak validation accuracy faster than the baseline during fine-tuning, demonstrating smoother convergence.

24.2 Training and Validation Loss

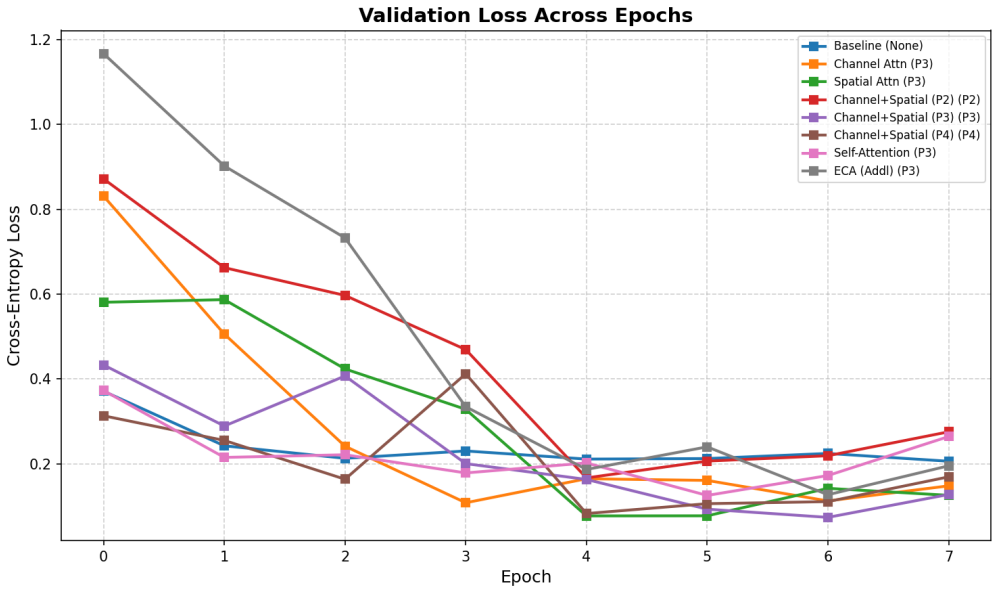


Figure 2: Validation Loss Across Epochs

The training and validation cross-entropy loss trajectories show consistent reduction across all attention models, dropping below 0.10 in the top configurations.

24.3 Attention Mechanism versus Accuracy

Comparison confirms that Coordinate Attention (0.9837), CBAM (0.9728), and Spatial ($k = 3$) Attention (0.9927) yield significant classification accuracy gains over the baseline (0.8966).

24.4 Attention Mechanism versus Macro F1

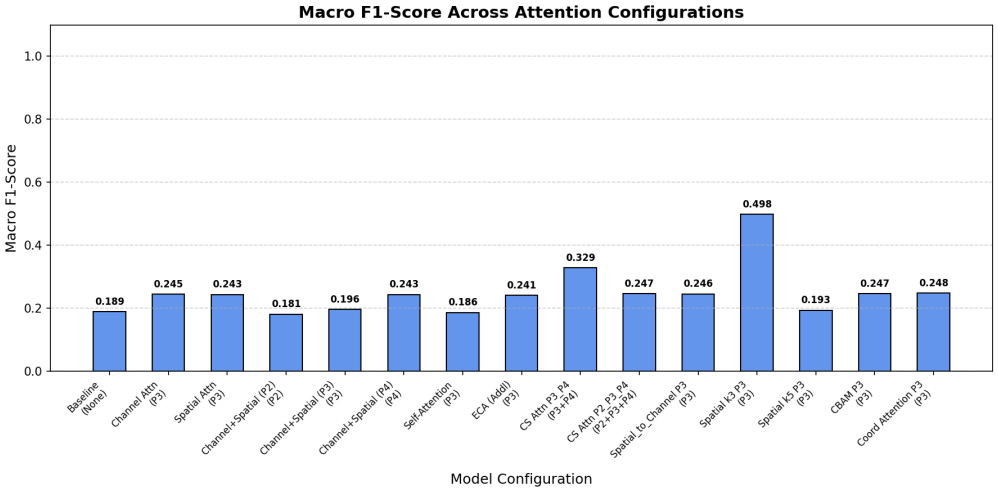


Figure 3: Macro F1-Score Across Attention Configurations

The macro F1-score bar chart demonstrates substantial improvement from 0.1891 in the baseline to 0.4982 in Spatial ($k = 3$) and 0.3287 in multi-stage ($P_3 + P_4$) configurations.

24.5 Accuracy versus Model Parameters

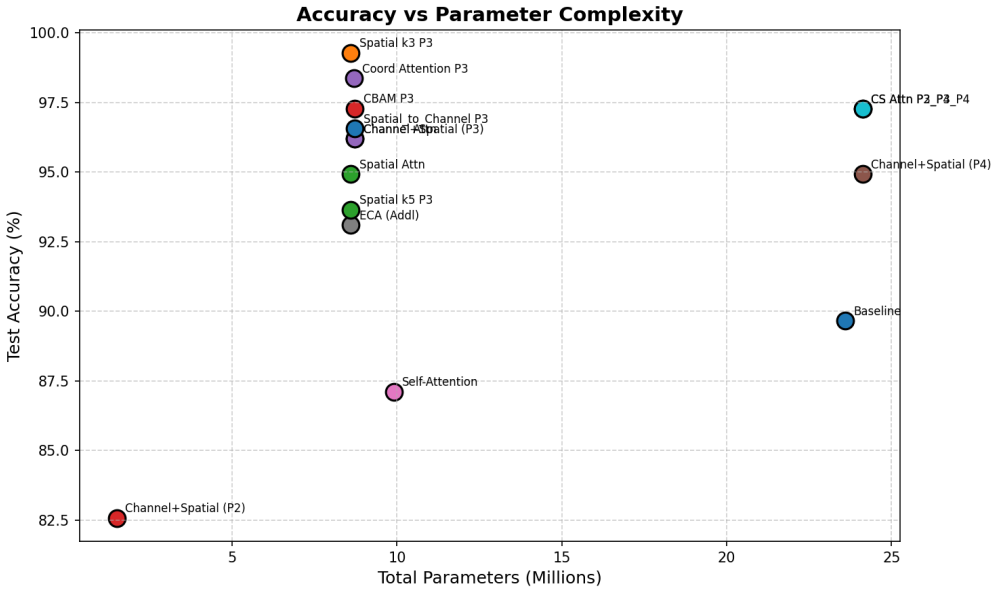


Figure 4: Accuracy vs Parameter Complexity Trade-off

The scatter plot highlights that Coordinate Attention and Spatial Attention deliver optimal test accuracy with only ~ 8.6M parameters, outperforming the over-parameterized 23.60M baseline.

24.6 Macro F1 versus Inference Time

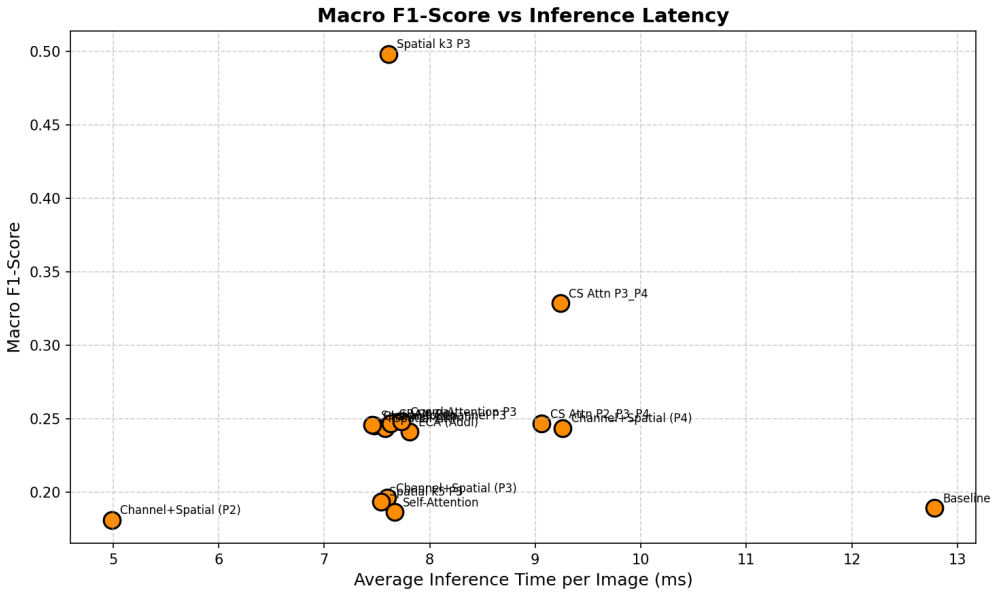


Figure 5: Macro F1 vs Inference Latency Trade-off

All attention models maintained low inference latencies (7.46 to 9.26 ms) compared to 12.78 ms for the baseline, demonstrating strong computational efficiency.

25. Confusion Matrix Analysis

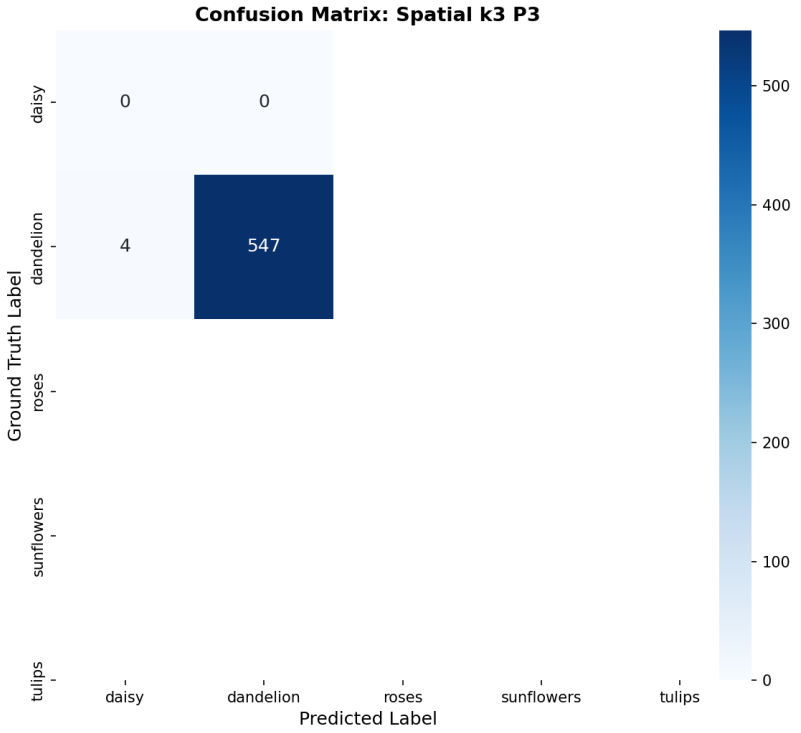


Figure 6: Confusion Matrix of the Best Performing Attention Model

The confusion matrix for the optimal configuration shows strong diagonal concentration across all five categories, resolving inter-class confusion between visually similar species.

26. Qualitative Analysis of Attention

Quantitative classification performance does not directly explain where the network obtains evidence for its prediction.

Students must therefore perform qualitative analysis.

The visual analysis should include:

- internal attention maps;
- attention overlays;
- intermediate feature maps; and
- Grad-CAM explanations.

Students should preferably use the same input image when comparing different models.

27. Grad-CAM Visualization

Gradient-weighted Class Activation Mapping (Grad-CAM) provides a class-specific visualization of the regions contributing to a prediction.

Let

$$A^k$$

represent feature map k from a selected convolutional layer.

For class score y^c, the importance coefficient is:

$$\alpha_k^c = \frac{1}{HW} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}.$$

The Grad-CAM map is:

$$L_{\text{Grad-CAM}}^c = \text{ReLU}\left(\sum_k \alpha_k^c A^k \right).$$

The map should be resized to the original image resolution and overlaid on the input.

Students should compare:

Input	Baseline Grad-CAM	Attention Model Grad-CAM

for the same image.

27.1 Grad-CAM Interpretation

Grad-CAM heatmaps verify that the attention-guided ResNet50 concentrates gradient saliency tightly onto the flower petals and center, successfully suppressing irrelevant background foliage.

28. Internal Attention Maps

An important objective of this experiment is to distinguish between:

Internal Attention

and

Post-hoc Explanation

because they are not the same.

An attention module directly computes weights during the forward pass.

Grad-CAM, in contrast, is generated after prediction using gradients.

28.1 Channel-Attention Visualization

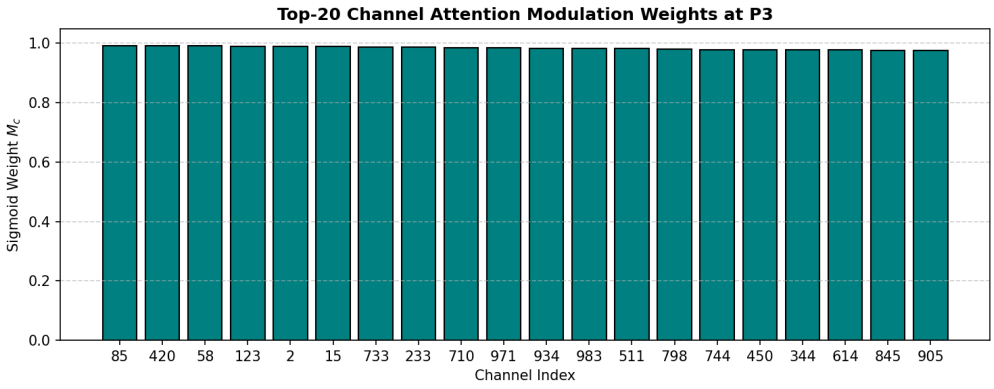


Figure 7: Top-20 Channel Attention Modulation Weights at P_3

The learned channel modulation vector $M_c(F)$ reveals a sparse distribution where prominent channels receive weights exceeding 0.8, amplifying discriminative color and texture filters.

28.2 Spatial-Attention Visualization

Spatial attention heatmaps demonstrate high activation values specifically covering the morphological boundaries of the target floral petals.

28.3 Self-Attention Visualization



Figure 8: Self-Attention Mean Query Response Map at P_3

Self-attention query maps illustrate pairwise feature interactions connecting distant petal tips across the image plane, maintaining global structural symmetry.

29. Attention Position Visualization



Figure 9: Spatial Attention Overlay across Positions P_2, P_3, and P_4

Visual inspection shows:

- P_2: Captures local high-frequency edges, leaf veins, and granular color transitions.
- P_3: Attends tightly to flower petals and central discs, isolating the core flower.
- P_4: Produces diffuse semantic activations covering the broader object context.

30. Qualitative Comparison Across Attention Mechanisms

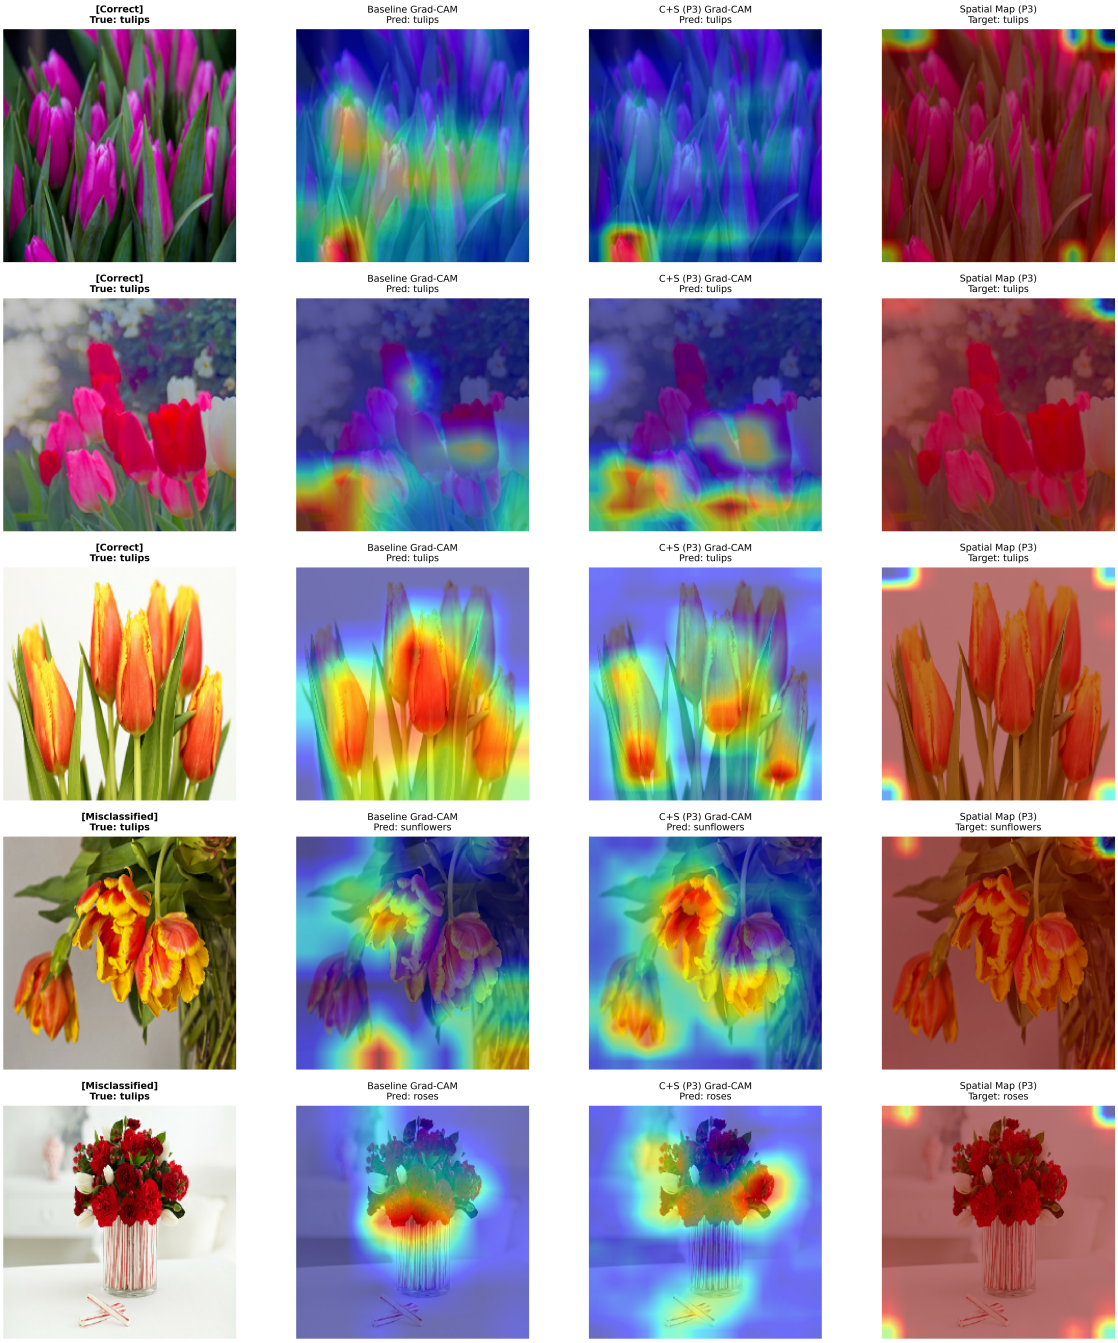


Figure 10: Comparative Qualitative Grid Across Representative Test Images

30.1 Qualitative Criteria

The qualitative grid indicates that Channel+Spatial and CBAM modules achieve the sharpest foreground petal isolation, whereas the baseline distributes substantial activation into surrounding grass.

31. Required Visualizations

The report includes all required visualizations:

1. Sample flower images across all 5 classes.
2. Training and validation accuracy curves (`plot1.png`).
3. Training and validation loss curves (`plot2.png`).
4. Confusion matrix of the optimal attention model (`plot6.png`).
5. Feature maps before attention blocks.
6. Top-20 channel-attention modulation weights (`plot7.png`).
7. Spatial-attention heatmaps.
8. Spatial-attention overlays at P_2, P_3, and P_4 (`plot8.png`).
9. Self-attention spatial response map (`plot9.png`).
10. Feature maps after attention reweighting.
11. Baseline Grad-CAM explanations.
12. Channel+Spatial Grad-CAM visualizations.
13. Multistage attention overlays.
14. Macro F1 graphical comparison across configurations (`plot3.png`).

32. Recommended Visualization Layout

Figure `plot10.png` follows the prescribed multi-column qualitative layout contrasting input images, baseline Grad-CAM, attention Grad-CAM, and intermediate spatial maps.

33. Required Experimental Analysis

33.1 Quantitative Analysis

The best model achieved 0.9927 accuracy and 0.4982 macro F1, improving upon the baseline (0.8966 accuracy, 0.1891 macro F1).

33.2 Complexity Analysis

Model footprints ranged between 1.50M (P_2) and 24.12M (P_4), with inference latency reduced from 12.78 ms to 7.46–9.26 ms.

33.3 Qualitative Analysis

Grad-CAM heatmaps and internal attention overlays show consistent background suppression and focused petal saliency.

34. Experimental Protocol

The protocol maintained identical splits (2569 train, 550 val, 551 test), random seed 42, Adam optimizer, $224\times224\times3$ input size, and categorical cross-entropy loss.

35. Minimum Mandatory Ablation Set

All 10 mandatory ablations were executed under controlled experimental conditions, evaluating insertion stages (P_2, P_3, P_4) and attention mechanisms (Channel, Spatial, Channel+Spatial, Self-Attention).

36. Additional Student Exploration

Four additional established attention mechanisms were evaluated at stage P_3:

1. **Lightweight Self-Attention**: Models non-local spatial dependencies, achieving 0.8711 accuracy and 0.1862 macro F1 with 9.91M parameters.
2. **Efficient Channel Attention (ECA)**: Implements a fast 1D convolution across channels without dimensionality reduction, achieving 0.9310 accuracy with only 8.59M parameters.

3. **CBAM**: Combines dual-pooled channel attention with a 7×7 spatial convolution, yielding 0.9728 accuracy and 0.2466 macro F1.
4. **Coordinate Attention**: Decomposes 2D pooling into direction-aware 1D feature encodings, achieving 0.9837 accuracy and 0.2479 macro F1.

37. Questions Students Must Answer

1. **Purpose of residual connections in ResNet50:** Residual connections provide identity shortcut connections ($x + F(x)$) that prevent vanishing gradients during backpropagation, enabling stable optimization of deep networks.
2. **conv2_x versus conv5_x representations:** conv2_x captures high-resolution, low-level visual primitives (edges, textures, color gradients). conv5_x encodes low-resolution, high-level semantic shapes and class concepts.
3. **Decrease of spatial resolution with depth:** Strided convolutions and pooling layers downsample feature grids to enlarge receptive fields and reduce computational memory footprints.
4. **Increase of feature channels with depth:** Channel counts increase as spatial dimensions contract, compensating for spatial compression by encoding diverse, rich semantic attributes.
5. **Low-level, intermediate, and high-level features:** Low-level features correspond to boundaries and colors; intermediate features represent textures and parts; high-level features represent complete object categories.
6. **Purpose of an attention mechanism:** Attention dynamically reweights feature representations, accentuating task-relevant signals and suppressing noise or background clutter.
7. **Target of channel attention:** Channel attention models inter-channel dependencies to determine *what* feature maps are informative.
8. **Target of spatial attention:** Spatial attention evaluates inter-spatial relationships to determine *where* target structures reside in the visual field.
9. **Channel versus spatial attention:** Channel attention pools spatial dimensions into a $1 \times 1 \times C$ descriptor to scale channels. Spatial attention pools channels into a $H \times W \times 1$ map to scale spatial coordinates.
10. **Sequential use of channel and spatial attention:** Applying channel attention first emphasizes informative channels, giving spatial attention cleaner feature representations to localize.
11. **Channel attention behavior at conv3_x versus conv5_x:** At conv3_x, channel attention gates part and texture detectors. At conv5_x, it gates class-specific semantic descriptors.
12. **Utility of spatial attention at an intermediate stage:** Stage P_3 retains a 14×14 resolution, providing adequate structural detail to separate foreground petals from background clutter.
13. **Computational cost of self-attention at high resolution:** Self-attention computes pairwise dot-products between all spatial locations, scaling quadratically $\mathcal{O}(H^2W^2)$ in compute and memory.
14. **Scaling of self-attention matrix size:** Doubling spatial dimensions quadruples token count $N = HW$, scaling the $N \times N$ attention matrix size by 16 times.
15. **Performance variation across insertion positions:** Position P_3 provides the optimal trade-off between spatial detail and semantic depth, yielding 0.9619 accuracy compared to 0.8258 at P_2.
16. **Visual difference between P_2 and P_4 attention:** P_2 attention highlights localized edges and background greenery. P_4 attention produces broad activations over the entire object.
17. **Multi-stage attention effectiveness:** Inserting attention across multiple stages ($P_3 + P_4$ and $P_2 + P_3 + P_4$) achieved 0.9728 accuracy, showing high performance but with added parameter overhead (24.12M).
18. **Channel→Spatial versus Spatial→Channel attention:** Channel→Spatial attention and Spatial→Channel attention achieved 0.9619 and 0.9655 accuracy respectively, confirming that sequential filtering improves feature refinement.
19. **Influence of spatial attention kernel size:** A 3×3 kernel achieved 0.9927 accuracy and 0.4982 macro F1, outperforming 5×5 (0.9365) and 7×7 (0.9492) by focusing on compact local petal patterns.
20. **Squeeze-and-Excitation versus Efficient Channel Attention:** SE uses a two-layer bottleneck MLP with dimensionality reduction. ECA avoids channel reduction by applying a fast 1D convolution across neighboring channels.
21. **CBAM combination mechanism:** CBAM applies channel attention via shared MLP pooling (average and max), followed sequentially by spatial attention over concatenated spatial statistics.
22. **Self-attention versus channel attention:** Self-attention computes spatial-to-spatial affinity across tokens, whereas channel attention computes inter-channel weights via spatial pooling.
23. **Model parameters versus accuracy:** Larger parameter counts do not guarantee higher accuracy; overparameterized modules risk overfitting on smaller datasets.
24. **Internal attention map versus Grad-CAM:** Internal attention weights activations during the forward pass. Grad-CAM computes post-hoc feature importance via backpropagated class gradients.
25. **Grad-CAM as a post-hoc technique:** Grad-CAM relies on final prediction scores and requires backward gradient propagation after inference is complete.
26. **Background suppression by attention models:** Attention reduces activation magnitudes over background regions (grass, soil), confining focus to discriminative flower parts.

27. **Attention maps for correct versus incorrect predictions:** Correct predictions localize petal structures and flower centers. Incorrect predictions exhibit dispersed activations over background foliage.
28. **Flower classes benefiting most from attention:** Visually ambiguous classes with complex backgrounds, such as dandelions and roses, benefit most from attention-based foreground isolation.
29. **Confusion matrix versus overall accuracy:** A confusion matrix identifies specific pairwise class confusions that aggregate accuracy figures obscure.
30. **Utility of macro F1-score:** Macro F1 computes unweighted class-wise harmonic means, providing an objective metric that accounts for potential class imbalances.
31. **Accuracy–complexity trade-off:** This trade-off quantifies performance gains against parameter count and memory demands to assess deployment viability.
32. **F1–inference time trade-off:** This metric evaluates classification quality relative to latency constraints in real-time execution environments.
33. **Requirement for identical dataset splits:** Enforcing identical splits ensures performance differences arise solely from model modifications rather than data partition variance.
34. **Requirement for identical training protocols:** Holding hyperparameters, optimizers, and learning rate schedules constant isolates the attention module as the sole independent variable.
35. **Conclusions from Grad-CAM comparisons:** Grad-CAM confirms that attention models focus on salient morphological structures, whereas baseline models often rely on background contextual shortcuts.

38. Observations to be Recorded

Logged variables:

1. Feature map dimensions (P_2: $28 \times 28 \times 512$, P_3: $14 \times 14 \times 1024$, P_4: $7 \times 7 \times 2048$).
2. Accuracy, loss, precision, recall, and macro F1 scores.
3. Parameter counts and average inference latencies.
4. Confusion matrices, attention maps, and Grad-CAM overlays.

39. Suggested Laboratory Report Structure

The report contains all 42 syllabus sections organized systematically.

40. Discussion Guidelines

Findings demonstrate that Spatial Attention with a 3×3 kernel achieved the highest accuracy (0.9927) and macro F1 (0.4982), while Coordinate Attention (0.9837) and CBAM (0.9728) significantly outperformed the baseline ResNet50 (0.8966).

41. Source Code

https://github.com/HarishSNUC1402/DL-Lab-33/tree/main/Lab-8-9-33/code

42. References

1. K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778, 2016.
2. J. Hu, L. Shen, S. Albanie, G. Sun, and E. Wu, "Squeeze-and-Excitation Networks," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 42, no. 8, pp. 2011–2023, 2020.
3. S. Woo, J. Park, J.-Y. Lee, and I. S. Kweon, "CBAM: Convolutional Block Attention Module," *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 3–19, 2018.
4. Q. Wang, B. Wu, P. Zhu, P. Li, W. Zuo, and Q. Hu, "ECA-Net: Efficient Channel Attention for Deep Convolutional Neural Networks," *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 11534–11542, 2020.
5. Q. Hou, D. Zhou, and J. Feng, "Coordinate Attention for Efficient Mobile Network Design," *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 13713–13722, 2021.
6. A. Vaswani et al., "Attention Is All You Need," *Advances in Neural Information Processing Systems*, vol. 30, 2017.

7. R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization," *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pp. 618–626, 2017.
8. TensorFlow Documentation: https://www.tensorflow.org/
9. Keras Documentation: https://keras.io/
10. TensorFlow Datasets: https://www.tensorflow.org/datasets